

Project SCC 2024

Report 1

Work done by:

João Lima nº 60350

Diogo Nunes nº 70502

Motivation

It was given to us the task to develop TuKano, an app inspired by TikTok, to leverage Microsoft Azure Cloud platform, in order to achieve better throughput on a regional scale. This report's objective is to identify in what ways the goal was achieved, and how managing Microsoft Azure Cloud platform has changed the base app given.

What was achieved

In our project, we implemented:

- Blob Storage, to save the blobs regarding the shorts posted;
- NoSQL Cosmos DB, to manage the users, shorts, likes and follows;
- PostgreSQL Cosmos DB, for an alternative storage (*);
- Redis Cache, for faster access to objects required.

To understand how these implementations impacted TuKano, we will look at different combinations of the previous and compare its metrics.

(*) - PostgreSQL Cosmos DB was fully coded in this project, but there was a connection error that was impossible for us to debug in the given project time, even with help from the practical classes teacher, Kevin Gallagher. We've decided to leave the metrics regarding the PostgreSQL Cosmos DB out of the report, but mention it in the implementations.

Impact on TuKano

We will be looking at these combinations:

1. Blob Storage + NoSQL Cosmos DB + Redis Cache
2. Blob Storage + NoSQL Cosmos DB

User_register

- With Redis Cache

http.codes.400:	400
http.codes.404:	200
http.downloaded_bytes:	15200
http.request_rate:	6/sec
http.requests:	600
http.response_time:	
min:	2

```

max: ..... 3446
mean: ..... 94
median: ..... 3
p95: ..... 210.6
p99: ..... 820.7
http.response_time.4xx:
min: ..... 2
max: ..... 3446
mean: ..... 94
median: ..... 3
p95: ..... 210.6
p99: ..... 820.7
http.responses: ..... 600
plugins.metrics-by-endpoint./tukano/rest/users/.codes.400: ..... 200
plugins.metrics-by-endpoint./tukano/rest/users/{{ userId }}?pwd={{ pwd }}.co... 200
plugins.metrics-by-endpoint./tukano/rest/users/{{ userId }}?pwd={{ pwd }}.cod... 200
plugins.metrics-by-endpoint.response_time./tukano/rest/users/:
min: ..... 2
max: ..... 356
mean: ..... 5
median: ..... 3
p95: ..... 4
p99: ..... 7
plugins.metrics-by-endpoint.response_time./tukano/rest/users/{{ userId }}?pwd={{ pwd }}:
min: ..... 2
max: ..... 10
mean: ..... 2.8
median: ..... 3
p95: ..... 4
p99: ..... 6
plugins.metrics-by-endpoint.response_time./tukano/rest/users/{{ userId }}?pwd={{ pwd }}:
min: ..... 181
max: ..... 3446
mean: ..... 274.3
median: ..... 198.4
p95: ..... 278.7
p99: ..... 2836.2
vusers.completed: ..... 200
vusers.created: ..... 200
vusers.created_by_name.TuKanoWholeUserFlow: ..... 200
vusers.failed: ..... 0
vusers.session_length:
min: ..... 192.8
max: ..... 3843.4
mean: ..... 287.9
median: ..... 206.5
p95: ..... 295.9
p99: ..... 2836.2

```

- Without Redis Cache

```
http.codes.400: ..... 400
http.codes.404: ..... 200
http.downloaded_bytes: ..... 15200
http.request_rate: ..... 6/sec
http.requests: ..... 600
http.response_time:
  min: ..... 1
  max: ..... 2627
  mean: ..... 82.3
  median: ..... 3
  p95: ..... 202.4
  p99: ..... 399.5
http.response_time.4xx:
  min: ..... 1
  max: ..... 2627
  mean: ..... 82.3
  median: ..... 3
  p95: ..... 202.4
  p99: ..... 399.5
http.responses: ..... 600
plugins.metrics-by-endpoint./tukano/rest/users/.codes.400: ..... 200
plugins.metrics-by-endpoint./tukano/rest/users/{{ userId }}?pwd={{ pwd }}.co... 200
plugins.metrics-by-endpoint./tukano/rest/users/{{ userId}}?pwd={{ pwd }}.cod... 200
plugins.metrics-by-endpoint.response_time./tukano/rest/users/:
  min: ..... 2
  max: ..... 302
  mean: ..... 4.7
  median: ..... 3
  p95: ..... 4
  p99: ..... 5
plugins.metrics-by-endpoint.response_time./tukano/rest/users/{{ userId }}?pwd={{ pwd }}:
  min: ..... 1
  max: ..... 15
  mean: ..... 2.6
  median: ..... 2
  p95: ..... 4
  p99: ..... 8.9
plugins.metrics-by-endpoint.response_time./tukano/rest/users/{{ userId}}?pwd={{ pwd }}:
  min: ..... 171
  max: ..... 2627
  mean: ..... 239.5
  median: ..... 194.4
  p95: ..... 252.2
  p99: ..... 1939.5
vusers.completed: ..... 200
```

```

vusers.created: ..... 200
vusers.created_by_name.TuKanoWholeUserFlow: ..... 200
vusers.failed: ..... 0
vusers.session_length:
  min: ..... 183.4
  max: ..... 2960.3
  mean: ..... 252.1
  median: ..... 202.4
  p95: ..... 262.5
  p99: ..... 1939.5

```

- Conclusion

In this case, Redis Cache does not improve response times and in some cases the latency increases, for example for the endpoints. This can be explained by the fact that in this case cache is only used for adding new members, not fulfilling its goal.

Upload_shorts

- With Redis Cache

```

http.codes.403: ..... 30
http.codes.404: ..... 30
http.downloaded_bytes: ..... 0
http.request_rate: ..... 6/sec
http.requests: ..... 60
http.response_time:
  min: ..... 3
  max: ..... 429
  mean: ..... 105.3
  median: ..... 175.9
  p95: ..... 202.4
  p99: ..... 210.6
http.response_time.4xx:
  min: ..... 3
  max: ..... 429
  mean: ..... 105.3
  median: ..... 175.9
  p95: ..... 202.4
  p99: ..... 210.6
http.responses: ..... 60
plugins.metrics-by-endpoint./tukano/rest/blobs/{{ blobUrl }}.codes.403: ..... 30
plugins.metrics-by-endpoint./tukano/rest/shorts/{{ userId }}?pwd={{ pwd }}.c... 30
plugins.metrics-by-endpoint.response_time./tukano/rest/blobs/{{ blobUrl }}:
  min: ..... 3

```

```

max: ..... 429
mean: ..... 18
median: ..... 4
p95: ..... 5
p99: ..... 5
plugins.metrics-by-endpoint.response_time./tukano/rest/shorts/{{ userId }}?pwd={{ pwd }}:
min: ..... 176
max: ..... 211
mean: ..... 192.5
median: ..... 190.6
p95: ..... 202.4
p99: ..... 206.5
vusers.completed: ..... 30
vusers.created: ..... 30
vusers.created_by_name.Upload short: ..... 30
vusers.failed: ..... 0
vusers.session_length:
min: ..... 188.4
max: ..... 671.7
mean: ..... 220
median: ..... 202.4
p95: ..... 228.2
p99: ..... 232.8

```

- Without Redis Cache

```

http.codes.403: ..... 30
http.codes.404: ..... 30
http.downloaded_bytes: ..... 0
http.request_rate: ..... 7/sec
http.requests: ..... 60
http.response_time:
min: ..... 2
max: ..... 266
mean: ..... 97.2
median: ..... 169
p95: ..... 198.4
p99: ..... 262.5
http.response_time.4xx:
min: ..... 2
max: ..... 266
mean: ..... 97.2
median: ..... 169
p95: ..... 198.4
p99: ..... 262.5
http.responses: ..... 60
plugins.metrics-by-endpoint./tukano/rest/blobs/{{ blobUrl }}.codes.403: ..... 30

```

```

plugins.metrics-by-endpoint./tukano/rest/shorts/{{ userId }}?pwd={{ pwd }}.c... 30
plugins.metrics-by-endpoint.response_time./tukano/rest/blobs/{{ blobUrl }}:
  min: ..... 2
  max: ..... 260
  mean: ..... 11.5
  median: ..... 3
  p95: ..... 4
  p99: ..... 4
plugins.metrics-by-endpoint.response_time./tukano/rest/shorts/{{ userId }}?pwd={{ pwd }}:
  min: ..... 170
  max: ..... 266
  mean: ..... 183
  median: ..... 179.5
  p95: ..... 198.4
  p99: ..... 198.4
vusers.completed: ..... 30
vusers.created: ..... 30
vusers.created_by_name.Upload short: ..... 30
vusers.failed: ..... 0
vusers.session_length:
  min: ..... 179.2
  max: ..... 547.3
  mean: ..... 202.8
  median: ..... 186.8
  p95: ..... 210.6
  p99: ..... 214.9

```

● Conclusion

Redis Cache does not have a significant role in Upload_shorts, even getting worse response times in some cases. Same as User_register, this is because a cache purpose is to retrieve wanted objects or information quicker, not create.

*This test had an error that it didn't recognize the path towards shorts.map and shorts.list that we could not solve, so it didn't upload the photos so it probably uploaded an empty file.

Realistic_flow

● With Redis Cache

```

errors.No shorts exist yet.: ..... 10
http.codes.404: ..... 20
http.downloaded_bytes: ..... 0
http.request_rate: ..... 4/sec
http.requests: ..... 20

```

```

http.response_time:
  min: ..... 170
  max: ..... 983
  mean: ..... 284
  median: ..... 190.6
  p95: ..... 757.6
  p99: ..... 757.6
http.response_time.4xx:
  min: ..... 170
  max: ..... 983
  mean: ..... 284
  median: ..... 190.6
  p95: ..... 757.6
  p99: ..... 757.6
http.responses: ..... 20
plugins.metrics-by-endpoint./tukano/rest/shorts/{{ shortId }}/{{ userId }}/l... 1
plugins.metrics-by-endpoint./tukano/rest/shorts/{{ userId }}/feed?pwd={{ pwd... 10
plugins.metrics-by-endpoint./tukano/rest/shorts/{{ userId }}/followers?pwd={... 1
plugins.metrics-by-endpoint./tukano/rest/shorts/{{ userId }}/shorts.codes.404: . 6
plugins.metrics-by-endpoint./tukano/rest/shorts/{{ userId1 }}/{{ userId2 }}/... 2
plugins.metrics-by-endpoint.response_time./tukano/rest/shorts/{{ shortId }}/{{ userId
  }}/likes?pwd={{ pwd }}:
  min: ..... 175
  max: ..... 175
  mean: ..... 175
  median: ..... 175.9
  p95: ..... 175.9
  p99: ..... 175.9
plugins.metrics-by-endpoint.response_time./tukano/rest/shorts/{{ userId }}/feed?pwd={{ pwd
  }}:
  min: ..... 170
  max: ..... 241
  mean: ..... 190.2
  median: ..... 186.8
  p95: ..... 198.4
  p99: ..... 198.4
plugins.metrics-by-endpoint.response_time./tukano/rest/shorts/{{ userId }}/followers?pwd={{
  pwd }}:
  min: ..... 983
  max: ..... 983
  mean: ..... 983
  median: ..... 982.6
  p95: ..... 982.6
  p99: ..... 982.6
plugins.metrics-by-endpoint.response_time./tukano/rest/shorts/{{ userId }}/shorts:
  min: ..... 287
  max: ..... 762
  mean: ..... 376.2

```



```

median: ..... 295.9
p95: ..... 314.2
p99: ..... 314.2
plugins.metrics-by-endpoint.response_time./tukano/rest/shorts/{{ userId1 }}/{{ userId2
}}/followers?pwd={{ pwd }}:
min: ..... 175
max: ..... 187
mean: ..... 181
median: ..... 175.9
p95: ..... 175.9
p99: ..... 175.9
vusers.completed: ..... 20
vusers.created: ..... 30
vusers.created_by_name.Download short: ..... 6
vusers.created_by_name.Follow user: ..... 2
vusers.created_by_name.Get Short Likes: ..... 4
vusers.created_by_name.Get User Follows: ..... 1
vusers.created_by_name.Get User's Shorts: ..... 6
vusers.created_by_name.Like short: ..... 1
vusers.created_by_name.View feed: ..... 10
vusers.failed: ..... 10
vusers.session_length:
min: ..... 178.5
max: ..... 1003.4
mean: ..... 291.6
median: ..... 194.4
p95: ..... 772.9
p99: ..... 772.9

```

- Without Redis Cache

```

errors.No shorts exist yet.: ..... 12
http.codes.404: ..... 18
http.downloaded_bytes: ..... 0
http.request_rate: ..... 4/sec
http.requests: ..... 18
http.response_time:
min: ..... 166
max: ..... 3177
mean: ..... 653.2
median: ..... 368.8
p95: ..... 2101.1
p99: ..... 2101.1
http.response_time.4xx:
min: ..... 166
max: ..... 3177

```

```

mean: ..... 653.2
median: ..... 368.8
p95: ..... 2101.1
p99: ..... 2101.1
http.responses: ..... 18
plugins.metrics-by-endpoint./tukano/rest/shorts/{{ shortId }}/{{ userId }}/l... 5
plugins.metrics-by-endpoint./tukano/rest/shorts/{{ userId }}/feed?pwd={{ pwd... 9
plugins.metrics-by-endpoint./tukano/rest/shorts/{{ userId }}/followers?pwd={... 1
plugins.metrics-by-endpoint./tukano/rest/shorts/{{ userId }}/shorts.codes.404: . 1
plugins.metrics-by-endpoint./tukano/rest/shorts/{{ userId1 }}/{{ userId2 }}/... 2
plugins.metrics-by-endpoint.response_time./tukano/rest/shorts/{{ shortId }}/{{ userId
}}/likes?pwd={{ pwd }}:
min: ..... 166
max: ..... 408
mean: ..... 275.4
median: ..... 247.2
p95: ..... 383.8
p99: ..... 383.8
plugins.metrics-by-endpoint.response_time./tukano/rest/shorts/{{ userId }}/feed?pwd={{ pwd
}}:
min: ..... 179
max: ..... 3177
mean: ..... 971.4
median: ..... 407.5
p95: ..... 2101.1
p99: ..... 2101.1
plugins.metrics-by-endpoint.response_time./tukano/rest/shorts/{{ userId }}/followers?pwd={{
pwd }}:
min: ..... 687
max: ..... 687
mean: ..... 687
median: ..... 685.5
p95: ..... 685.5
p99: ..... 685.5
plugins.metrics-by-endpoint.response_time./tukano/rest/shorts/{{ userId }}/shorts:
min: ..... 371
max: ..... 371
mean: ..... 371
median: ..... 368.8
p95: ..... 368.8
p99: ..... 368.8
plugins.metrics-by-endpoint.response_time./tukano/rest/shorts/{{ userId1 }}/{{ userId2
}}/followers?pwd={{ pwd }}:
min: ..... 187
max: ..... 392
mean: ..... 289.5
median: ..... 186.8
p95: ..... 186.8

```

```

p99: ..... 186.8
vusers.completed: ..... 18
vusers.created: ..... 30
vusers.created_by_name.Download short: ..... 9
vusers.created_by_name.Follow user: ..... 2
vusers.created_by_name.Get Short Likes: ..... 3
vusers.created_by_name.Get User Follows: ..... 1
vusers.created_by_name.Get User's Shorts: ..... 1
vusers.created_by_name.Like short: ..... 5
vusers.created_by_name.View feed: ..... 9
vusers.failed: ..... 12
vusers.session_length:
  min: ..... 169.5
  max: ..... 3189.6
  mean: ..... 662.6
  median: ..... 391.6
  p95: ..... 2143.5
  p99: ..... 2143.5

```

- Conclusion

In this case, the use of Redis Cache was very valuable:

- Response Times and Consistency improved:

With Redis Cache, response times were significantly lower (653.2 ms to 284 ms) and latency was better (P95, P99). For this topic, Redis Cache helped by retrieving user data more quickly, reducing the database workload and improving latency.

- There were less 404 Errors:

This can be an error in the code, maybe the shorts weren't being correctly deleted from the cache so sometimes they had been already eliminated before and it still had access to them, minimizing some of the 404 Errors (Not Found). Nonetheless, consistency improved.

Overall, there is a significant improvement in the metrics by using Redis Cache, this being the best example from the three to show what changed.

Final Conclusion

Same as mentioned in the beginning, the objective was to compare the results of using PostgreSQL Cosmos DB (which is fully implemented) with NoSQL Cosmos DB but unfortunately we had an error regarding the PostgreSQL Cosmos DB and this was not achieved.

As we expected, the cache increased the time of POST operations but decreased the time of GET operations, which in the context of Tukano is better since in a real life scenario, GET operations are more used. Concluding, **Redis Cache had a positive effect on TuKano.**